
Reading Without Looking: A Position-Resolved Profile of Confidence and Correctness in a Small Document OCR Model

Adya Srivastava
srivastavadya@gmail.com

Abstract

Confidence scores from document OCR models are used to decide which predictions a human should check. We audit that signal in a small (≈ 19.6 M parameter) encoder-decoder OCR model trained from scratch with no prior exposure to any Indic script, across three independently trained seeds. We first establish what the model can do: on held-out Hindi images its grapheme character error rate is 0.985, and on blank white images it is 0.949. The model does not read the held-out images, which are drawn from GlotOCR Bench and rendered by an independent pipeline, and text-bearing and blank input are not distinguishable at three seeds ($t = 0.257$, $p = 0.822$, 2 df). We then localise the failure. Under teacher forcing the ground-truth token is never the model’s argmax at generation position 0 (0 of 180 sequences), and the geometric mean probability assigned to the correct first symbol is 2.2×10^{-11} on text-bearing images and 6.3×10^{-11} on blank: more than eight orders of magnitude below what a text-only n -gram model fitted on the same training corpus, which never sees an image, assigns to that same symbol. This is confirmed by a scale-free statistic immune to softmax-sharpness objections: the median rank of the correct symbol among 367 candidates is 70.5, with 35.0% of sequences ranking it below position 100. Self-generated max-softmax confidence at that position is nonetheless approximately 0.90. From position 2 to 39 the model recovers to between -0.15 and -0.48 , comparable to a position-matched 4- to 5-gram grapheme prior and near-identical with and without the image; at positions 0, 1 and beyond 39 it falls far below every n -gram baseline. Deleting the encoder changes the output distribution substantially (mean KL 1.075 nats, argmax flips at 8.2% of decoding steps) but leaves its peak near unity: 96.7% of the divergence sits at those flipped positions while agreeing positions retain a mean confidence of 0.998. Finally, the confidence signal is not globally broken. On synthetic lines drawn entirely from the training manifest, where the model attains 18.3% line accuracy, confidence discriminates correct from incorrect transcriptions at AUROC 0.838 while remaining badly miscalibrated at a mean of 0.994 — a discrimination result we can state only for memorised text, since this evaluation set has no held-out arm; on the held-out images, where accuracy is zero, it carries no usable correctness signal. We read this as a model whose certainty tracks the linguistic plausibility of what it generates rather than the state of its own perception. Because visual dependence in transcription is concentrated at the start of the sequence, a confidence score averaged uniformly over generated tokens dilutes the one position that carries it; confidence measures for transcription should weight early positions rather than averaging uniformly. We do not exclude the possibility that this instrument never acquired a visual representation that transfers to the evaluation renderer, and we state what would separate that from visual bypass.

1 Introduction

Document OCR is often the first automated step in processing multilingual records, and production pipelines route low-confidence predictions to human review. The design assumes that when a model is unsure, it says so.

That assumption is unreliable for large vision–language models. Recent work establishes *visual ungroundedness*: models produce fluent, confident, sometimes correct output driven entirely by language priors, with the image contributing nothing, and existing confidence estimators cannot detect this [Khanmohammadi et al., 2026]. Related work describes the same hazard as perception bypass or visual shortcut [Xu, 2026], and multilingual OCR benchmarks report that failure on unfamiliar scripts is not silent: models hallucinate fluent text in familiar scripts rather than abstaining [Kargaran et al., 2026].

We take that as established and ask a narrower question that autoregressive transcription makes answerable: **where in the generated sequence does visual dependence live, and can its absence be localised without training a probe?**

The question has a structure here that short-answer visual question answering does not offer. When a decoder transcribes a line token by token, every position after the first can draw on two sources: the image, through cross-attention, and the preceding tokens, through the language prior over the training script. These are confounded. Position 0 is not. At position 0 there is no prefix, so the only *input-dependent* source of information is the image; whatever the decoder contributes is a fixed marginal over sequence-initial symbols that cannot vary with the input. Position 0 therefore isolates visual dependence in a way no later position does, provided the comparison is made against the model’s own distribution at that position rather than against a uniform prior. Measuring it requires only ground-truth text and two forward passes: no auxiliary labels, no trained detector, no second model at measurement time. *Interpreting* the measurement does require a reference point from a model with demonstrated reading ability, and Section 10 reports what we can and cannot say without one.

We build a small encoder–decoder OCR model from scratch rather than fine-tuning an existing checkpoint, so that “this model has never seen this script” is a claim we can make rather than assume. We then run four interventions on fixed checkpoints across three seeds: encoder output deletion, blank and unfamiliar-script input, teacher-forced ground-truth likelihood, and a positional decomposition of that likelihood.

Contributions. (1) We establish by direct measurement that the instrument does not read its held-out images (grapheme CER 0.985 text-bearing against 0.949 blank, a difference that does not survive seed clustering), and treat this as a finding to be explained rather than a caveat to be buried; every subsequent claim is scoped by it. (2) We give a position-resolved grounding diagnostic for autoregressive transcription, motivated by the unconfounded status of position 0 and requiring only ground-truth text and two forward passes. (3) We decompose encoder ablation into distributional *shape* (KL divergence, large, 96–99% concentrated at argmax flips) and *peak height* (max-softmax, 0.998 at agreeing positions), showing that the quantity deployed as confidence is the one component of the output distribution that visual ablation leaves near ceiling. (4) We identify a regime in which the same confidence signal does discriminate correctness (AUROC 0.838 at 18.3% accuracy) while remaining uncalibrated, giving a within-instrument contrast rather than a purely negative result; we are explicit that this evaluation set has no held-out arm, so the result is stated for memorised text and not offered as evidence of generalisation. (5) We show that *at least* 16.9–20.4% of non-exact predictions from off-the-shelf OCR engines on Indic document images are Unicode encoding variants rather than misreadings, on the two engines with adequate sample size and with a large unreviewed remainder, a precondition for any grounding claim measured against exact-match references.

What this paper does not claim. We did not discover that confidence fails to track visual grounding; that is prior work [Khanmohammadi et al., 2026]. We make no claim about production-scale OCR systems. And we do not yet have a positive control, a model with demonstrated reading ability on which the same diagnostic returns a different answer; Section 10 states what that costs.

2 Related Work

Visual bypass and language priors. That a multimodal model can answer without using the image originates in visual question answering and has been rediscovered repeatedly. Contemporary work treats the existence of such perception bypass as a documented hazard rather than a novel observation [Xu, 2026]. A blind language model ignoring image evidence has been shown to outperform prior art on some image–text retrieval benchmarks [Lin et al., 2024].

Positional structure of visual dependence. That visual reliance is not uniform across generation positions is established. M3ID [Favero et al., 2024] shows that as more tokens are generated the reliance on the visual prompt decreases, names this *conditioning dilution*, and builds a decoder around a per-position comparison of the output distribution with and without the image. First Logit Boosting [Ha et al., 2026] reports that the first generated token carries the strongest visual evidence and reuses its logit as a persistent anchor against the same decay. Both operate on large vision–language models that read their inputs, and both treat the positional structure as a means to a decoding method. We study the boundary case instead: a model with no usable visual representation, where the gradient those works describe has an empty source. Our question is not whether visual reliance declines with position, which is known, but what the profile looks like when there is no visual reliance to decline from, and whether the confidence signal tracks it.

Confidence estimation under ungroundedness. BICR [Khanmohammadi et al., 2026] is the closest prior work. It establishes that existing confidence estimators cannot separate grounded from ungrounded predictions, and proposes a probe trained on real-image hidden states, regularised by a ranking loss against a blacked-out view, evaluated across five large vision–language models. Our contribution differs in being positional rather than learned, training-free, and aimed at sequence transcription rather than short-answer generation. We do not offer a better estimator; we offer a cheaper diagnosis of a specific structural cause.

Abstention and probing in OCR. Latent-state and attention probes have been trained for OCR-error abstention, including visual-focused probes motivated by the informativeness of attention to visual tokens [Yao et al., 2025]. HALP predicts hallucination risk before generation from internal representations, reaching up to 0.93 AUROC on several models [Kogilathota et al., 2026]. Both require supervision; the diagnostic here does not.

Multilingual OCR and grapheme-aware evaluation. GlotOCR Bench finds that models hallucinate fluent text in familiar scripts rather than abstaining, with even the best model transcribing under 8% of low-resource-tier sentences below 5% CER [Kargaran et al., 2026]. A Devanagari stress-test benchmark finds that synthetic renders badly overstate quality, that nine of ten evaluated systems collapse on real scans, and that English OCR strength does not predict Indic performance [Singh, 2026]. *grapheme-kit* extends evaluation metrics from Unicode code points to grapheme clusters and shows via an OCR case study that grapheme-level metrics evaluate complex scripts more faithfully [Nisfer et al., 2026]; Section 4 applies that argument rather than claiming it.

3 The Instrument

A small encoder–decoder vision–language model, approximately 19.6M parameters. A Vision Transformer encoder trained from random initialisation, with no pretrained weights, produces patch-level visual features. An autoregressive decoder cross-attends to those features and emits one grapheme cluster per step. The output vocabulary is defined over grapheme clusters, the visually atomic units of the script, rather than byte-pair-encoded subwords, so one output token corresponds to one decision about what was read. Vocabulary size is $|V| \approx 367$, built from the run’s own manifest with a minimum grapheme frequency of 5. Three seeds were trained independently under identical hyperparameters, a design adopted after a single-seed result failed to replicate (Appendix A).

Where the images come from. Both training and evaluation text originate in the Hindi portion of GlotOCR Bench [Kargaran et al., 2026], but the images do not. Training images are produced by our own renderer from GlotOCR Hindi ground-truth strings. Evaluation images are GlotOCR Bench’s own renders of Hindi sentences, in its two variants: a clean rendering on a white canvas

and a degraded rendering that applies paper texture, ink wear, geometric distortion, downsampling and compression artifacts. **They are renders, not photographs or scans of physical documents**, and we do not describe them as scans anywhere in this paper. The distribution shift the instrument faces at evaluation is therefore between two rendering pipelines rather than between synthetic and photographed text, which is a smaller gap than a scan corpus would present. Throughout, we write *text-bearing* for evaluation images containing rendered text, in contrast to *blank* white images. We avoid the word “real” for our own data entirely, since none of it is photographed.

The training pipeline includes a renderer with a glyph-frequency dial supporting three target distributions (natural, flattened, inverted), achieving total-variation distance ≤ 0.08 from target in all three. As a measurement and control instrument it worked; as a data generator it did not. Line accuracy across the three conditions is $18.3\% \pm 5.1$ (natural), $0.33\% \pm 0.6$ (flattened) and $0.67\% \pm 0.6$ (inverted), seed means over three seeds. With the dependent variable at the floor in two of three conditions, a fixed-effects decomposition of exposure against intrinsic script complexity is uninterpretable, and we withhold it (Section 10).

4 What Counts as an Error

Indic orthography admits multiple valid representations of visually identical text, so exact-match scoring conflates encoding variation with misreading. Every claim below is measured against ground-truth strings, so the reference must be established first. We distinguish three categories. **Tier 1 (encoding)**: deterministic byte-level variation with no difference in meaning or rendering, covering Unicode normalisation forms, zero-width joiner and non-joiner variants around conjuncts, danda versus period, script-native versus Latin digits, and anusvara versus homorganic nasal under a fixed sandhi rule; computed by a deterministic normalisation function with no judgment involved. **Tier 2 (phonetic)**: whether two spellings represent the same intended word, assessed by canonicalisation through ISO 15919 rather than a hand-enumerated rule list, validated against 38 hand-constructed pairs (23 positive, 15 negative or boundary). **Human residual**: genuine misreading, dropped diacritics, reading-order breaks, hallucinated or repeated content, assessed by grapheme-cluster alignment rather than code-point alignment.

We apply this taxonomy to three off-the-shelf OCR engines on the same image corpus. The scope matters and we state it before the numbers: the engine study pools prediction rows across four Indic languages (Hindi, Bengali, Santhali, Kashmiri) and both plain and degraded image variants, and it scores full strings rather than the grapheme-level alignment used for the model in Section 5. It establishes that exact-match scoring misrepresents Indic OCR error rates in general. It is not a Devanagari-specific measurement and it is not on the same footing as the model-side results.

Table 1: Tier 1 encoding variants among predictions that do not exactly match the reference, pooled over four Indic languages and both image variants; n counts prediction rows, not documents. The final column is the quantity of interest and is a lower bound, since a large share of non-exact predictions is unreviewed. PaddleOCR is shown for completeness only; $n = 10$ cannot support an engine-level claim.

Engine	n	Exact	Tier 1	Tier 1 among non-exact
Tesseract	180	15.6%	17.2%	20.4%
Surya	222	46.8%	9.0%	16.9%
PaddleOCR	10	0%	10.0%	10.0%

Between roughly one sixth and one fifth of what naive exact-match scoring calls an OCR error on these documents is an encoding variant, not a misreading. Two caveats belong with that number. First, a large share of non-exact predictions remains unreviewed (60.0% for Tesseract, 44.1% for Surya, 90.0% for PaddleOCR), so the Tier 1 shares are lower bounds on the encoding-artifact rate rather than final estimates. Second, Tier 2 resolves *zero* additional cases on all three engines: once Tier 1 normalisation is applied, phonetic spelling variation contributes nothing measurable on this corpus. We report the Tier 2 machinery because it was validated and applied, and its null result is itself informative about which kinds of orthographic ambiguity actually arise in Indic OCR output. All model-side results below use Tier-1-normalised references and grapheme-cluster alignment.

5 The Instrument Does Not Read Its Held-Out Images

We report this before any confidence result, because it scopes everything that follows.

Table 2: Grapheme CER by condition and seed, after Tier-1 normalisation with grapheme-cluster segmentation and Levenshtein distance. Lower is better; 1.0 corresponds approximately to a prediction sharing nothing with the reference.

Condition	Seed 0	Seed 1	Seed 2	Pooled
Hindi, held-out ($n = 60/\text{seed}$)	1.009	0.984	0.963	0.985
Blank white ($n = 60/\text{seed}$)	0.830	0.896	1.122	0.949
Ol Chiki, unseen ($n = 100/\text{seed}$)	0.974	0.906	1.029	0.970
Perso-Arabic, unseen ($n = 20/\text{seed}$)	0.958	0.906	1.110	0.991

Three things follow.

First, **the model does not read its held-out images**. A CER near 1.0 is not partial reading; it is a transcription sharing almost nothing with the reference.

Second, **the nominal blank advantage is not robust and we do not build on it**. The sign is not stable across seeds: blank CER is lower at seeds 0 and 1 (+0.180 and +0.088) and higher at seed 2 (−0.159). A Wilcoxon signed-rank test on the 60 paired images is significant unclustered ($p = 0.025$) but not once clustered by seed ($t = 0.257$, $p = 0.822$). Generated length explains 23–24% of CER variance in both conditions, and on exact length-matched pairs ($n = 54$) the gap is +0.039 with $p = 0.197$. Note that blank outputs are not shorter than text-bearing ones (mean grapheme length 31.4 against 28.0), so the mechanism is repetition rather than truncation: Section 6.2 shows blank generations collapsing onto a handful of repeated strings. **The defensible claim is that text-bearing and blank CER are statistically indistinguishable**, which is sufficient for everything that follows.

Third, **the unseen-script conditions are not a distinct regime**. Ol Chiki (0.970) and Perso-Arabic (0.991) sit inside the same band as in-distribution Hindi (0.985). The model fails on held-out images of its own training script as thoroughly as on scripts it has never seen.

The consequence for framing is direct. We cannot claim a dissociation between reading ability and confidence, because on these images there is no reading ability to dissociate from. What we can claim, and what follows, is a mechanistic account of how a model in this state produces near-ceiling confidence, and a measurement that localises it.

Bypass or domain shift? A reader may object that this instrument was trained on images from our own renderer and evaluated on images from GlotOCR Bench’s, so it may have no usable visual representation for the evaluation distribution at all. That is domain shift rather than bypass, and the two are different failures. We note that the gap is narrower than the objection usually assumes: both sides are rendered text from the same source corpus, in the same script, at the same canonical line height, differing in rendering stack, font and degradation pipeline rather than in medium. Our own evidence cuts both ways. Against bypass: zeroing the encoder changes the argmax at 8.2% of decoding steps (Section 7), so the encoder does contribute something. For bypass: the position profile is near-identical on text-bearing and blank input at every bucket (Section 8), and confidence on two scripts the model has never seen is indistinguishable from in-distribution confidence (Table 3), neither of which domain shift alone predicts. We also ran the two measurements that would separate them: Gaussian-noise images and patch-scrambled real images, testing whether the model responds to visual variation of any kind rather than only to the specific rendering gap between training and evaluation. Pooled whole-sequence mean $\log p(\text{GT})$ is −1.756 on noise and −1.730 on patch-scrambled input, against −1.783 text-bearing and −1.751 blank — all four conditions within 0.05 nats of one another, and the position-bucket profile for both new conditions reproduces the same shape as text-bearing and blank: collapse at positions 0 (−23.7 noise, −24.1 scrambled) and 1 (−8.9, −7.9), recovery through positions 2–39, and a second collapse beyond position 40. Pure noise carries no structure a domain-shift account would expect the model to process differently from a familiar-but-shifted rendering, and patch-scrambled input destroys spatial layout while preserving local texture statistics; if the model retained any domain-specific visual competence, these two conditions should diverge from text-bearing input more than blank does. They do not. This is ev-




Figure 1: Text-bearing, blank, Gaussian-noise, and patch-scrambled input compared. **Left:** whole-sequence mean $\log p(\text{GT})$, per-seed points with the pooled mean (dashed). **Right:** $\log p(\text{GT})$ at position 0 specifically, same layout. All four conditions cluster tightly on both measures; noise and patch-scrambled input are indistinguishable from text-bearing and blank rather than falling between them, which is the pattern a residual, domain-specific visual competence would predict and does not appear here.

idence for bypass over domain shift, though we note it is evidence, not proof: an instrument that never learned to use images under any condition is consistent with both a total bypass failure and a domain-shift failure severe enough to leave no residual sensitivity to any visual input at all.

We also state the training budget plainly, because it bears on this and a reader will reach it in Appendix E otherwise: 2,538 rendered line crops, a single unrecorded font, no augmentation, 5,000 steps at a constant learning rate, and the final checkpoint taken without validation-based selection. That is a small budget, and a reader is entitled to conclude the model may simply never have acquired a visual representation that transfers across renderers. We regard that as an open alternative rather than a settled one. What it does not explain on its own is why confidence is equally high on scripts the model has never seen and on blank pages, nor why the position profile is unchanged when the image is removed entirely.

6 Confidence Under Absent and Unfamiliar Input

Table 3: Mean generation confidence (max-softmax of the selected token, averaged over the generated sequence). Seed columns give the mean over 60 images with the standard deviation across those images; the pooled column is over $180 = 60 \times 3$ observations on 60 shared images, not 180 independent images (Section 6.4).

Condition	Seed 0	Seed 1	Seed 2	Pooled
Hindi, text-bearing ($n=60$)	0.9824 ± 0.0198	0.9897 ± 0.0085	0.9899 ± 0.0119	0.9873
Blank white ($n=60$)	0.9814 ± 0.0125	0.9933 ± 0.0084	0.9824 ± 0.0154	0.9857
OI Chiki ($n=100$)	0.9848 ± 0.0139	0.9847 ± 0.0117	0.9876 ± 0.0120	0.9857
Perso-Arabic ($n=20$)	0.9901 ± 0.0072	0.9894 ± 0.0109	0.9887 ± 0.0083	0.9894

All four conditions sit between 0.985 and 0.990. The spread across condition means (0.0037) is smaller than the seed-to-seed spread of those means. Which seed trained the model accounts for more variation in reported confidence than whether the input carries text, an unfamiliar script, or a blank page.

We note a statistical caution rather than leaving it to a reader. An equivalence test at a smallest-effect-of-interest of $\delta = 0.05$ is weak on this data: every condition lies within a 0.005-wide band, so equivalence within $\delta = 0.05$ is close to guaranteed by the measurement scale rather than by the finding. The bound would be meaningful only if calibrated against the spread a grounded model




Figure 2: Per-image mean confidence by condition, faceted by seed, with the seed mean and standard deviation overlaid. Note the truncated vertical axis, $[0.92, 1.00]$: the point of the figure is the narrowness of the band, not its position. The four distributions occupy the same band in every seed, and seed-to-seed movement of the band exceeds any between-condition separation within a seed.

exhibits over the same conditions. We therefore report the equivalence test as descriptive and rest no claim on it.

6.1 What the model emits on unfamiliar scripts

On all 360 held-out images spanning both unseen scripts and all three seeds, the model produces zero characters of the target script, substituting text composed 93–100% (per-condition mean) of Devanagari. Manual inspection shows fluent, grammatical Hindi unrelated to the source image. This is the small-scale counterpart of what GlotOCR Bench reports at scale [Kargaran et al., 2026].

6.2 Output degeneracy

Confidence equivalence across conditions is interesting only if the outputs compared are non-trivial. They are not equally so.

Table 4: Output diversity within each seed, $n = 60$ images. At seed 1 the median pairwise grapheme edit distance between outputs on *different* blank images is 0.

Condition, seed	Unique outputs / 60	Mean pairwise edit	Identical-pair rate
Hindi, seed 0	24	26.4	12%
Hindi, seed 1	30	29.4	5%
Hindi, seed 2	22	22.7	18%
Blank, seed 0	3	7.5	34%
Blank, seed 1	3	7.2	76%
Blank, seed 2	8	26.7	27%

First-token identity shows the same pattern. On blank input at seed 0 the decoder begins every one of 60 generations with the same grapheme; at seed 1, two distinct first graphemes cover all 60, one accounting for 86.7%. On text-bearing Hindi the corresponding counts are 10, 9 and 10 distinct first graphemes, with modal frequencies between 40% and 72%.

We draw the honest conclusion: part of the “confidence is unchanged on blank input” result is that the decoder emits a near-constant, highly probable string on blank input. That is a real property of the model and consistent with the paper’s thesis, but it is a weaker mechanism than fluent varied hallucination and we do not describe it as the latter. Text-bearing input is also far from diverse: 22 to 30 unique outputs across 60 distinct images.

6.3 Mismatched pairing rules out the degeneracy confound

Blank is a degenerate input by construction, so “blank scores similarly to text-bearing” is weaker evidence than a design that removes degeneracy entirely: score each ground-truth string against a *different, real* image, under a fixed derangement of the 60-image pool, so every comparison image is a genuine text-bearing photograph and none is a repeat of its own match. Mean teacher-forced $\log p(\text{GT})$ under this mismatch is -1.552 (seed 0), -1.702 (seed 1), -1.871 (seed 2); medians -1.133 , -1.326 , -1.436 . These sit inside the same range as the correctly-paired text-bearing condition reported in Section 8 (-1.783) and the blank condition (-1.751), not below them. The position-bucket breakdown is the same shape too: pooled mean $\log p(\text{GT})$ under mismatch is -23.80 at position 0, -7.82 at position 1, recovering to -0.47 , -0.18 and -0.24 across positions 2–39, and collapsing again to -5.90 beyond position 40 – indistinguishable in shape from Table 6’s text-bearing and blank rows. Scoring a caption against a real photograph of an entirely different page produces the same likelihood profile as scoring it against its own correct image or against a blank one. The degeneracy explanation in Section 6.2 accounted for why blank output looks artificially close to text-bearing output; it cannot account for this, since the comparison images here are genuine, varied photographs, not a collapsed handful of repeated strings. This closes the alternative that the paper’s central result is an artefact of what the decoder does on blank pages specifically.

6.4 Variance structure

Because the 180 observations per condition are 60 images under three seeds, pooled standard deviations mix image variance, seed variance and their interaction. Decomposed, confidence on text-bearing Hindi is 9% image / 7% seed / 84% residual, and on blank 0% / 18% / 82%. Ground-truth likelihood behaves oppositely at 84–85% image / $\approx 2\%$ seed / $\approx 14\%$ residual. Which image is presented explains 9% of confidence variance on text-bearing input and none of it on blank, while explaining most of the variance in a quantity that genuinely depends on the input. This is independent evidence that the confidence signal is not responding to input content, and it means pooled standard errors in Table 3 should be read as seed-clustered. Seed-clustered bootstrap intervals over 2,000 resamples give, for confidence on text-bearing Hindi, image 9.1% [0.0, 22.4], seed 7.1% [0.2, 17.7] and residual 83.9% [67.9, 96.6]; on blank, image 0.0% [0.0, 0.0], seed 18.0% [11.0, 27.7] and residual 82.0% [72.3, 89.0]. For ground-truth likelihood the image share is 84.1% [69.7, 91.1] on text-bearing and 85.0% [70.8, 90.9] on blank.

Two caveats. Greedy decoding is deterministic, so there is one observation per image–seed cell and the interaction is not separable from noise; what we label residual *is* the image $\times$ seed interaction. And the image share for confidence on text-bearing input has an interval that includes zero, so the 9% versus 0% contrast is suggestive and not a test. What the intervals do support is the difference in kind between the two quantities: the image share for ground-truth likelihood is bounded well away from zero, and for confidence it is not.

7 Mechanism I: Ablation Moves the Distribution’s Shape, Not Its Peak

We replace the encoder’s output with a zero tensor of identical shape immediately before the decoder’s cross-attention consumes it, and compare against the unmodified model on the same text-bearing images. Divergences are reported as $\text{KL}(p_{\text{full}} \parallel p_{\text{zero}})$ throughout.

Two generation procedures are involved and we separate them, because the distinction decides what the decomposition means. For the per-step divergence, top-1 agreement and their agree-versus-flip split, the zeroed model is *teacher-forced on the full model’s greedy token prefix* at every step, so both distributions are conditioned on an identical context and sequence divergence cannot confound the comparison. The context is therefore the full model’s own generation rather than the reference transcription. The confidence delta in the first column is different: it compares the full model’s greedy run against an *independent* zeroed greedy run, which is the quantity a deployed system would report. The two columns are not measurements of the same procedure and should not be read as such.

We give the per-seed rows before the pooled figure deliberately. The seed-level confidence deltas span 0.034 and change sign. **The pooled -0.0030 is cancellation, not a stable small effect, and**

Table 5: Encoder zeroing, $KL(p_{\text{full}} \parallel p_{\text{zero}})$ throughout. Image-level aggregates (left) and a step-level decomposition of the KL into positions where the argmax agrees between the full and zeroed model and positions where it flips (right). The two aggregations are not interchangeable and do not reconstruct one another arithmetically: the image-level columns average per-image means, weighting short and long sequences equally, while the step-level columns average over all 4,990 decoding steps. Top-1 agreement is correspondingly a per-image mean (12.1% disagreement pooled) and the flip rate is per step (8.2% pooled).

Seed	Image-level			Step-level KL decomposition			
	Δconf	KL	Top-1 agr.	Flip rate	KL agree	KL flip	Flip share
0	-0.0234	1.639	0.857	10.2%	0.050	10.62	96.0%
1	+0.0035	0.440	0.950	4.9%	0.017	8.92	96.3%
2	+0.0109	1.147	0.831	13.4%	0.010	6.48	99.0%
Pooled	-0.0030	1.075	0.879	8.2%	0.027	8.87	96.7%




Figure 3: **Left:** share of total ablation KL contributed by agreeing versus argmax-flipping decoding steps, per seed and pooled. The bars encode share of KL; the percentages in the legend are share of *steps*. Flipped steps are 8.2% of steps but carry 96.7% of the divergence. **Right:** mean top-1 confidence at agreeing and flipped positions, with the full and zeroed encoder. Agreeing positions are pinned at 0.998 under both; flipped positions remain high but not pinned.

the confidence consequence of encoder ablation is unmeasured at three seeds rather than measured to be zero.

Pooled mean confidence, split the same way: with the full encoder, 0.998 at agreeing positions and 0.893 at flipped positions; with the encoder zeroed, 0.998 and 0.825. The statement this supports, and which we make no stronger:

At the 88% of positions where deleting the encoder does not change which token wins, the distribution is essentially unchanged (KL 0.010–0.050) and the peak stays at 0.998. At the 12% where the encoder does change the winner, the distribution changes enormously (KL 6.5–10.6) and the model remains highly confident in the new winner (0.89 with the encoder, 0.83 without). The encoder can change *which* token holds the peak. It does not change *that there is a near-unit peak*.

Since deployed confidence is the peak height, and peak height is close to invariant under the intervention that removes visual information, confidence is structurally rather than incidentally blind to this class of failure. The invariance is strong but not total: confidence at flipped positions (0.89) is meaningfully below the agreeing-position value (0.998).

Two distinct interventions. This section zeroes the encoder output. Section 8 uses blank white images, where the encoder still runs and produces features of blankness. Zeroing asks whether the

cross-attention pathway carries usable signal at all; a blank image asks whether the model distinguishes content from its absence. We keep them labelled separately.

Cross-attention contribution norms. The behavioural claim above, that the encoder contributes something but not much, has a direct architectural measurement: $\|\text{cross-attn output}\|/\|\text{residual stream}\|$ per decoder layer, averaged over positions, computed separately on text-bearing and blank input across all three seeds and all five decoder layers. The pattern is consistent across every seed: the cross-attention pathway’s contribution to the residual stream is small in absolute terms (0.10–0.13 across layers and conditions) and is *higher on text-bearing input than on blank at every layer, in every seed* (e.g. seed 0, layer 0: 0.1131 blank versus 0.1202 text-bearing; seed 1, layer 2: 0.1183 versus 0.1427). The gap is modest, ranging roughly 0.01–0.02 in absolute norm depending on layer and seed, but its direction is uniform across all 15 (seed, layer) cells. This is independent, architecture-level evidence that the encoder pathway is not inert: it carries measurably more signal when there is something in the image to encode, even though that additional signal does not translate into correct output at position 0 or into a lower self-reported confidence when it is absent. Full per-layer, per-seed values are in Appendix C.

8 Mechanism II: Teacher-Forced Likelihood and the Position Profile

Max-softmax over a model’s own greedily generated tokens is upward-biased by construction. We therefore also compute teacher-forced log-likelihood of the ground-truth sequence, conditioning on true preceding tokens and recording the probability assigned to the true next token regardless of whether the model would have chosen it. Whole-sequence means are -1.783 (SD 1.274, $n = 180$) on text-bearing input and -1.751 (SD 1.280, $n = 180$) on blank, with mean predictive entropy 0.0207 and 0.0251 respectively.

The two measures disagree in sign, and we spin neither. Ground-truth likelihood is marginally higher on blank, contrary to what grounded confidence predicts. Predictive entropy is also higher on blank, meaning slightly less certain, which is exactly what grounded confidence predicts. Both differences are negligible against their dispersions: 0.032 against SDs near 1.27, and 0.004 against SDs near 0.017. Neither measure detects a difference between text-bearing and blank, and the nominal signs should not be interpreted. The two estimators are also not independent: both derive from the same softmax, differing in whether the conditioning prefix is self-generated or ground truth.

Where the ground truth wins. The ground-truth token is the model’s argmax at 90.4% of positions on text-bearing input and 90.1% on blank, with a mean peak probability of 0.9983 and 0.9982 at those positions, consistent with the reported entropy. Broken out by position: **0% at position 0**, on 0 of 180 sequences, under both conditions; 92.6% (text-bearing) versus 92.3% (blank) thereafter. That split is the shape of the whole result.

Table 6: Mean $\log p(\text{GT})$ by generation position, against n -gram grapheme language model references fitted on 2,491 of the 2,538 training lines (rows matching an evaluation string removed, add- α smoothing with $\alpha = 0.01$; note that the instrument itself was *not* held out this way, see Appendix E) and evaluated *within the same position buckets*, so model and reference are position-matched. The n -gram models never see an image. Token counts are over the 60 evaluation strings; instrument figures pool the same strings across three seeds. Uniform over $|V| \approx 367$ is -5.91 .

Positions	n tok.	Instrument		Position-matched n -gram reference				
		Text-bearing	Blank	1-gram	2-gram	3-gram	4-gram	5-gram
0	60	-24.54	-23.48	-5.92	-4.73	-4.73	-4.73	-4.73
1	60	-8.12	-8.63	-5.08	-1.38	-1.14	-1.14	-1.14
2–9	480	-0.48	-0.41	-4.11	-2.25	-0.88	-0.38	-0.31
10–19	598	-0.15	-0.16	-4.18	-2.34	-1.06	-0.45	-0.26
20–39	851	-0.27	-0.28	-4.13	-2.24	-0.98	-0.43	-0.20
40+	348	-6.12	-6.28	-4.17	-2.28	-0.99	-0.40	-0.18

Five observations, in order of importance.




Figure 4: The dissociation. **Top:** mean $\log p(\text{GT})$ against generation position, text-bearing and blank overlaid. The horizontal reference lines drawn here are *whole-tail* n -gram means (trigram -0.99 , 4-gram -0.44 , 5-gram -0.26) computed over all positions ≥ 1 ; Table 6 gives the position-matched references that the argument in this section uses, and those differ per bucket. Uniform is -5.91 . The inset zooms positions 2–39 to separate the 4-gram and 5-gram references, and the right-most point aggregates positions 40 and beyond and is a categorical bin rather than a position on the continuous axis. The vertical axis of the top panel is symmetric-log and therefore non-uniform: the interval from -25 to -10 occupies comparable space to the interval from -1 to 0 , which compresses the region containing the position-0 collapse. **Bottom:** self-generated max-softmax confidence at the same positions, axis clipped to $[0.5, 1.0]$. Correctness collapses at positions 0–1 while confidence does not, and both curves are near-identical between text-bearing and blank input.

1. Position 0 is far below a text-only baseline, not merely uncertain. The geometric mean probability assigned to the correct first symbol is 2.2×10^{-11} on text-bearing images and 6.3×10^{-11} on blank. The model is not undecided at position 0; it is near-deterministically committed to a specific wrong symbol, identically whether the image contains text or nothing. We report the geometric mean because the arithmetic mean of $p(\text{GT})$ is dominated by a small number of less-severe cases and is not representative of the distribution; Appendix B gives the per-seed figures and the histogram.

We deliberately do not anchor this on the uniform value of 2.725×10^{-3} . Self-generated max-softmax at position 0 is 0.902 (next paragraph), so the position-0 distribution is nowhere near uniform, and in a distribution that peaked almost any non-argmax symbol falls orders of magnitude below $1/|V|$. A comparison against uniform would therefore partly measure the sharpness of the softmax rather than a disfavouring of the truth specifically.

The position-matched n -gram references in Table 6 supply a null that is not subject to this objection, because both sides of the comparison are $p(\text{GT})$. A grapheme n -gram model fitted on the same training corpus, which never sees an image, assigns the correct first symbol -4.73 nats at position 0 (identical for orders two through five, since with no prefix they all reduce to the marginal over sequence-initial symbols). The instrument assigns it -24.54 . That is a gap of 19.8 nats, or more than eight orders of magnitude, *below a text-only baseline*. The model is not merely failing to use the image at position 0; it is placing the correct first symbol far lower than a model that has only ever seen text. We regard this, rather than the comparison against uniform, as the defensible form of the claim.

One further fact makes this harder to explain away rather than easier. The instrument was trained on the full manifest with no held-out split, and 19 of the 60 evaluation strings appear verbatim as training lines (Appendix E), while the n -gram references were fitted with the exactly-matching rows removed. The asymmetry is real but modest, and we state it at its true size: the reference is deduplicated rather than held out, since the other 41 evaluation strings are resampled from the same pool and their contexts remain in the fit, and dropping 47 of 2,538 rows moves a position-0 marginal by under two percent. What survives is the direction: the comparison at position 0 favours the instrument, and a model that has seen a third of these strings verbatim during training still never places the correct first symbol at the argmax and still assigns it roughly 20 nats less probability than the reference.

Two scale-free quantities sharpen the claim further, and both are now measured. Over the same 180 (image, reference) pairs, the median rank of the ground-truth symbol among the 367 candidates is 70.5 pooled (seed medians 60.5, 59.5, 95.0), IQR [33.0, 136.2], with 35.0% of sequences ranking the correct symbol below position 100 in the model’s own distribution. This is a scale-free statistic, immune to the softmax-sharpness objection raised above: a model that places the truth at a median rank of 70 out of 367 is not merely confident about the wrong symbol, it is actively disavowing the correct one relative to hundreds of alternatives. Mean predictive entropy at position 0 is 0.333 nats pooled, low enough to confirm the distribution is sharply peaked, consistent with the near-ceiling confidence reported next. The probability assigned to an arbitrary non-argmax symbol is also measured (Appendix B) and remains orders of magnitude above $p(\text{GT})$, confirming the collapse is specific to the correct symbol rather than an artefact of general softmax sharpness. The memorisation question for Section 9 is resolved, not outstanding: the synthetic evaluation set has no held-out arm, so the split cannot be run, and Section 9’s AUROC is reported as an in-training-manifest figure accordingly.

2. Confidence at that position is near ceiling. Self-generated max-softmax at position 0 is 0.902 pooled on text-bearing input and 0.895 on blank. Position 0 is therefore the point of maximal confidence and minimal correctness, at the one generation step where vision is the only possible information source. Per seed, text-bearing: 0.884 / 0.906 / 0.916; blank: 0.9999 / 0.898 / 0.787. The blank column is not stable across seeds and its pooled value should not be quoted alone.

3. Position 1 is also far below the text-only baseline. At -8.12 and -8.63 the second generated symbol sits 7.0 nats (text-bearing) and 7.5 nats (blank) below the best position-matched n -gram reference for that position (-1.14). The failure is not confined to a single step; recovery begins at position 2.

4. In the middle of the sequence the model is comparable to a high-order grapheme prior. Against position-matched references (Table 6), positions 2–9 (-0.48) sit between the trigram (-0.88) and 4-gram (-0.38) references; positions 10–19 (-0.15) are slightly *better* than the 5-gram reference (-0.26); positions 20–39 (-0.27) sit between the 4-gram and 5-gram references. So mid-sequence the instrument is in the same band as a 4- to 5-gram grapheme prior, occasionally above it, and it gets there without reference to the image, since the blank column matches the text-bearing column at every bucket.

Two properties of this comparison make the mid-sequence agreement weaker evidence than it looks, and we would rather state them than have them found. First, the two sides are not held out symmetrically: the n -gram references drop every manifest row matching an evaluation string, while the instrument was trained on the full manifest, so at positions where the instrument matches or exceeds the 5-gram reference it is doing so with an advantage the reference does not have. The -0.15 at positions 10–19, the one bucket where the instrument beats the 5-gram model, is exactly where that advantage would show. We therefore do not claim the instrument outperforms a high-order grapheme prior; we claim only that its mid-sequence likelihood sits in that band under a comparison it wins by construction.

Second, the training text is drawn from the same GlotOCR Hindi string pool that the evaluation strings come from, resampled into synthetic pages, and the n -gram models are fitted on that pool. A 5-gram grapheme model reaching -0.18 (roughly 0.84 per token) on about 2,500 lines implies that contexts recur heavily, which resampling produces by construction. The 5-gram reference is therefore closer to a memorisation baseline than to a generic language prior. This cuts in the direction




Figure 5: Per-position-bucket KL(model || 5-gram) (left) and argmax agreement rate between the model and the 5-gram reference (right), pooled bars with per-seed points overlaid. Both measures move together: low KL and high agreement in the recovered mid-sequence region, both collapsing sharply at positions 0, 1, and 40+ — exactly where the position profile itself changes shape.

of our claim rather than against it, since matching a memorising model mid-sequence is a stronger indictment than matching a generic one, but the comparison should be read with that in mind. It also means the 18.3% synthetic line accuracy in Section 9 likely *overstates* genuine reading ability, since the evaluation text is drawn from the same resampled pool as the training text; we report it as a regime marker rather than as a capability estimate.

A distributional comparison, replicated across all three seeds, supports the same picture. Per-position KL between the model and the 5-gram reference, together with the rate at which their argmaxes coincide, pooled over $n = 180$: at positions 2–9, 10–19 and 20–39, mean KL is 0.32, 0.27 and 0.28 nats and argmax agreement is 93.5%, 91.2% and 90.7% — the model and the 5-gram reference are not just close in mean log-likelihood, they select the same token at the overwhelming majority of mid-sequence positions, consistently across every seed (per-seed agreement at positions 2–9 ranges 92.9–93.8%). At positions 0, 1 and 40+, where the earlier observations report the instrument falling far below the n -gram baselines, KL rises sharply (4.98, 1.28 and 1.92 nats pooled) and argmax agreement collapses (4.4%, 62.8%, 61.5%). The two measures move together exactly where the position profile changes shape, replicated at every seed, which is what a genuine 4-to-5-gram account predicts and a coincidental mean-matching would not. We regard the exclusive claim as supported.

5. Outside that band the model is far worse than a text-only baseline. The position-matched view makes the shape of the failure precise, and it is not the shape a low-order- n -gram account predicts. At position 0 the instrument is 19.8 nats below the best n -gram reference; at position 1 it is 7.0 nats below; and beyond position 39 it is 5.9 nats below the 5-gram reference, on 348 evaluation tokens. The n -gram baselines are essentially flat across buckets (the 5-gram reference moves only from -0.31 to -0.18), so the instrument’s profile is not tracking any property of the text. It matches a text-only prior in the middle of the sequence and collapses far beneath one at both ends. We place limited weight on the 40+ bucket, since sequences reaching that length are a selected minority, but the position-0 and position-1 gaps rest on all 180 sequences.

9 When the Confidence Signal Does Work

A purely negative result would leave open the possibility that this model’s confidence head is simply broken. It is not, and the contrast is informative.

Before the numbers, a scoping note that governs all of them, and it is stronger than we initially stated. The instrument was trained on the full rendered manifest with no held-out split (Appendix E), and we checked directly whether the synthetic evaluation lines behind this section’s numbers are held out from that manifest: they are not, for any of them. All 300 evaluation instances (100 ground-truth strings $\times$ 3 seeds) are sampled from the training manifest itself, so there is no non-matching subset to

Table 7: Confidence as a predictor of correctness across evaluation regimes; n pools 100 or 60 images over three seeds. AUROC is against line-level exact correctness except in the final row, which uses a median split on grapheme CER because exact accuracy is identically zero. On the flattened and inverted sets at most one line per seed is exactly correct, so AUROC there is computed on one or zero positive cases and carries no information; we mark it uninformative rather than reporting the value.

Evaluation set	n	Line acc.	AUROC	Spearman (conf, CER)
Synthetic, natural exposure	300	0.183	0.838	−0.40
Synthetic, flattened / inverted	300 / 300	0.003 / 0.007	uninformative	≈ 0
Held-out Hindi images	180	0.000	undefined	−0.15
Held-out Hindi, median CER split	180	—	0.57	—




Figure 6: **Left:** confidence as a ranker of correctness in two regimes. The synthetic curve is against line-level exact correctness; the held-out curve is against a median split on grapheme CER, because exact accuracy there is identically zero and the AUROC against it is undefined (Table 7). The two curves therefore answer slightly different questions and should not be read as a like-for-like comparison. **Right:** reliability diagram on synthetic natural lines, equal-mass confidence deciles. Confidence is near unity in every decile while accuracy is not, so the model discriminates without being calibrated.

compare against — the memorisation-versus-correctness split we describe below was attempted and could not be run, because the design has no held-out arm. The 18.3% line accuracy and the AUROC that follows are therefore **in-training-manifest figures only**. They do not show, and cannot show with this evaluation set, that confidence tracks correctness on text the model has not seen. We use them here only as a *regime marker*: this identifies a setting in which the model produces some correct output on training-distribution text, so that the confidence signal’s behaviour in that setting can be compared against the setting where it produces none, not as evidence about generalisation.

The alternative we cannot exclude is therefore not merely plausible but the default reading: if the lines the model transcribes exactly are disproportionately ones it memorised, a confidence signal that ranks them highly may be tracking memorisation rather than correctness, and the two hypotheses produce the same AUROC on this evaluation set by construction. Testing which is true requires a held-out evaluation split that does not currently exist for the synthetic condition, and building one is future work, not a rerun of an existing probe.

On these synthetic lines, where the model gets 18.3% of lines exactly right, confidence ranks correct transcriptions above incorrect ones at AUROC 0.838. That is a working discrimination signal (per seed 0.763, 0.924, 0.832; bootstrap 95% interval [0.771, 0.902]), and it is nonetheless badly calibrated: mean confidence in that regime is 0.994 against 18.3% accuracy, giving an expected calibration error of 0.811 over ten equal-mass bins. Discrimination and calibration come apart cleanly here, and only the former survives. On the held-out images, where accuracy is zero, a graded substitute gives 0.57, near chance. Mean confidence across the three synthetic exposure conditions also varies substantially, at 0.994 (natural), 0.600 (flattened) and 0.461 (inverted), tracking the accuracy drop in those conditions.

Taking this together with Sections 5–8, the reading best supported by the evidence is that **the model’s confidence responds to properties of the text it is generating, not to whether the image supports that text**. It falls when the target text has unfamiliar statistics. It does not fall when the image is blank, unreadable, or in a script the model has never seen, because in those cases the decoder still produces fluent, familiar Devanagari. We offer this as an interpretation consistent with what we measured, not as a demonstrated mechanism.

10 Limitations and Future Work

No positive control; this is the central limitation. We have no model with demonstrated Devanagari reading ability on which the position-0 measurement returns a different answer. Without that contrast the measurement has no established dynamic range, and a reader cannot distinguish a working instrument from a broken one on this evidence alone. We attempted Surya, one of the engines in Table 1 and an architecture in the same encoder–decoder family as our instrument, and it is not usable for this purpose: the deployed inference package serves a llama.cpp-backed VLM with no exposed softmax and no separable encoder tensor to zero, and the older checkpoint matching our intended architecture is incompatible with current loading code (mismatched attention head dimensions, additional mixture-of-experts parameters with no principled way to drop them). We declined to force the load with a mismatched-size override, since that would silently produce meaningless weights rather than a real result. We attempted a second candidate, PaddleOCR’s own recognition backbone (already in Table 1 and confirmed to read Devanagari), and it fails for a more fundamental reason than Surya did: the installed recognition model (devanagari_PP-OCRv5_mobile_rec) is CTC-based, not autoregressive. Its SVTR component is a visual encoder over the feature map, not a text decoder; output is a softmax over the charset at each spatial column, collapsed by CTC decoding, with no beginning-of-sequence token and no $p(g_i | g_{<i}, \text{image})$ to teacher-force. There is no generation loop to extract per-step logits from, and no natural mapping from CTC’s column-wise output to our grapheme-cluster vocabulary. Reading column-0 softmax as a substitute for position-0 teacher-forced likelihood would not measure the same quantity, so we do not report one. Broader benchmarks of real-scan Devanagari competence [Singh, 2026] identify stronger candidates, though the top-scoring systems there are closed and do not expose per-step likelihoods. A viable positive control remains the single most valuable open experiment, and both candidates attempted so far have failed for architectural reasons rather than for lack of trying: one serves an inference stack with no exposed distribution at all, the other has no autoregressive distribution to expose in the first place.

What remains. Mismatched pairing (Section 6.3), cross-attention contribution norms (Section 7), noise and patch-scrambled input (Section 5), and the n -gram KL/argmax-agreement decomposition (Section 8), replicated across all three seeds, are now measured and no longer listed here. The memorisation question for Section 9 is also resolved, though not in the way we originally planned: we checked whether the synthetic evaluation set has a non-matching subset to split against, and it does not, so the split cannot be run on this evaluation set. What remains is therefore not a re-analysis but a new evaluation set: synthetic lines rendered from text that provably does not appear in the training manifest, generated and held out before any training occurs, would let Section 9’s AUROC be tested for generalisation rather than reported as an in-training-manifest figure. This requires new rendered data, not new model training or new GPU time. A positive control on a model with demonstrated reading ability remains the single item in this list that could not be completed within available compute; the probe implementations needed for it are written and released with the code, so running them is a matter of finding a compatible model, not further development.

No held-out training split. The instrument was trained on the entire rendered manifest, with no evaluation strings withheld; 19 of the 60 strings used for the teacher-forced analyses appear verbatim as training lines, and the remainder are resampled from the same source pool. This is a design fault rather than an analysis choice, and a redesign should hold out the evaluation strings before training rather than after. Its consequences are not uniform across the paper and we have tried to mark each one where it falls. It inflates the synthetic line accuracy in Section 9, which we therefore use only as a regime marker, and it removes the possibility of testing whether the AUROC reported there measures confidence tracking correctness or confidence tracking memorisation: we checked directly, and all 300 evaluation instances are drawn from the training manifest with no non-matching subset to compare against, so the split we originally proposed cannot be run on this

evaluation set. Section 9’s AUROC should be read as an in-training-manifest figure, not as evidence about generalisation. It biases the mid-sequence comparison in Section 8 in the instrument’s favour, since the n -gram references were held out and the instrument was not, so we do not claim the instrument beats a high-order prior. It does not weaken the position-0 and position-1 results, which are the paper’s central measurements: contamination can only have helped the instrument there, and it still fails at floor. And it is orthogonal to the encoder-ablation results in Section 7, which compare the same model against itself on the same images.

Scale and generality. One model of approximately 19.6M parameters, one training script for the mechanistic experiments. We make no claim about production OCR systems. A separate attempt to use TrOCR-base-printed as an external comparison found mean confidence 0.42 on its own in-domain English text, far below this instrument’s near-ceiling regime; that is an operating-regime mismatch rather than a failed replication, and it does not serve as a control.

Code and data. The instrument, the renderer, all probe implementations and the committed probe outputs from which every number in this paper is computed are available at <https://github.com/Adya6714/vlm-ocr-eval>.

Metric choice. Line-level exact match is uninformative here and is used only in Section 9, where the comparison is between regimes rather than between conditions. Grapheme CER is what Section 5 reports.

The exposure-versus-complexity question is unanswered. We believe the fix is to abandon distributional-target synthesis in favour of subsampling real, naturally learnable sentences to bias exposure, accepting weaker distributional control in exchange for text that remains learnable.

Novelty relative to prior work. Two things in this paper are not new and we do not present them as such. That a multimodal model’s confidence is insensitive to visual ablation is established [Khanmohammadi et al., 2026]. That visual reliance declines across generation positions, and that the first generated token carries the most visual evidence, are established too [Favero et al., 2024, Ha et al., 2026], on models that read their inputs and as the basis for decoding methods rather than as objects of study.

What we add is the boundary case and one decomposition. Prior work characterises the *gradient* of visual reliance in grounded models; we characterise the profile when there is no visual reliance for it to decline from, and find that position 0 is not merely the least language-supported position but one where the ground-truth symbol is never the argmax across 180 sequences, while self-reported confidence there is at its maximum. Separately, the shape-versus-peak decomposition in Section 7 is a sharper statement than “confidence fails to detect ungroundedness”: ablation moves *which* token holds the peak without moving *that* there is a near-unit peak, and the deployed scalar is the component it leaves untouched. The within-instrument regime contrast in Section 9 is what keeps this from being a purely negative result.

11 Discussion and Conclusion

Three implications follow from the position profile, if it holds under a positive control. Visual dependence in autoregressive transcription is concentrated at the start of the sequence: positions 2 to 39 are recoverable from the prefix alone at 62–86% per-token likelihood, in the band of a position-matched 4- to 5-gram grapheme prior, and near-identically so with no image at all. Any confidence score averaged uniformly over the generated sequence therefore averages the one vision-dependent position against many that are not, diluting exactly the signal a gating system needs; confidence measures for transcription should weight early positions. The diagnosis is also cheap: two forward passes and ground-truth text give the position profile, with no probe training and no auxiliary labels. Interpreting it against a reference model is a separate requirement, and one we have not met (Section 10). Where probe-based approaches [Khanmohammadi et al., 2026, Yao et al., 2025, Kogilathota et al., 2026] detect ungroundedness by learning it, the position-0 test asks a structural question with a known answer under grounding. And the failure mode is a training pathology that standard monitoring would not catch: a model in this state converges, produces fluent in-script output, and reports mean confidence of 0.99.

We audited a small, fully exposure-controlled document OCR model and found that it does not read its held-out images. We localised the failure to the start of the generated sequence, where the ground-truth symbol is never the model’s argmax across 180 sequences and receives a geometric mean probability of order 10^{-11} , more than eight orders of magnitude below what a text-only n -gram model fitted on the same corpus assigns it, identically with and without visual content, while the model’s own selected symbol at that position receives 0.90. We showed that deleting the encoder alters the output distribution’s shape substantially while leaving its peak near unity, so that the quantity read as confidence is close to invariant under precisely the intervention that removes visual information. And we showed that the same confidence signal does discriminate correctness in a regime where the model reads, locating the failure in the relationship between confidence and vision rather than in the confidence head itself.

References

- Alessandro Favero, Luca Zancato, Matthew Trager, Siddharth Choudhary, Pramuditha Perera, Alessandro Achille, Ashwin Swaminathan, and Stefano Soatto. Multi-modal hallucination control by visual information grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14303–14312, 2024.
- Jiwoo Ha, Jongwoo Baek, and Jinhyun So. First logit boosting: Visual grounding method to mitigate object hallucination in large vision-language models. *arXiv preprint arXiv:2604.00455*, 2026.
- Amir Hossein Kargaran, Nafiseh Nikeghbal, Jana Diesner, François Yvon, and Hinrich Schütze. GlotOCR bench: OCR models still struggle beyond a handful of unicode scripts. *arXiv preprint arXiv:2604.12978*, 2026. Also in the ICML 2026 virtual programme.
- Reza Khanmohammadi, Erfan Miah, Simerjot Kaur, Charese H. Smiley, Ivan Brugere, Kundan Thind, and Mohammad M. Ghassemi. Grounded or guessing? LVLm confidence estimation via blind-image contrastive ranking. *arXiv preprint arXiv:2605.10893*, 2026.
- Sai Akhil Kogilathota, Sripadha Vallabha E G, Luzhe Sun, and Jiawei Zhou. HALP: Detecting hallucinations in vision-language models without generating a single token. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6067–6085, Rabat, Morocco, 2026. Association for Computational Linguistics.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 29914–29934. PMLR, 2024.
- Izzath Nisfer, Ashini Kavindya, Ovindu Atukorala, Purushoth Velayuthan, and Menan Velayuthan. grapheme-kit: Grapheme-level metrics and text processing for multilingual NLP. *arXiv preprint arXiv:2607.22456*, 2026.
- Aditya Pratap Singh. Can OCR-VLMs read devanagari? a stress-test benchmark and post-correction study. *arXiv preprint arXiv:2606.29213*, 2026.
- Zekun Xu. When does a video-language model stop watching? reward strength controls the formation and reversal of visual shortcuts in multimodal RLVR. *arXiv preprint arXiv:2606.22043*, 2026.
- Jihan Yao, Achin Kulshrestha, Nathalie Rauschmayr, Reed Roberts, Banghua Zhu, Yulia Tsvetkov, and Federico Tombari. Reading between the lines: Abstaining from VLM-generated OCR errors via latent representation probes. *arXiv preprint arXiv:2511.19806*, 2025.

A A Retracted Single-Seed Result

An earlier analysis using only seed 0 reported a statistically significant confidence difference on the Perso-Arabic condition relative to in-distribution Hindi ($|z| = 2.54$). It did not replicate: at seed 1, Hindi confidence (0.9897) exceeded Perso-Arabic (0.9894), reversing the original direction. We retract the claim in full and report it as the failure mode that motivated the three-seed design.

B Position-0 Likelihood Distribution

Per-seed geometric means of $p(\text{GT})$ at position 0 are 2.05×10^{-10} / 1.35×10^{-11} / 3.88×10^{-12} (text-bearing) and 9.32×10^{-10} / 3.13×10^{-11} / 8.70×10^{-12} (blank). The effect is present independently at every seed and spans two orders of magnitude across seeds, so per-seed values rather than the pooled geometric mean should be used when comparing against another model.

A geometric mean can be driven by a minority of extreme cases, so we report the distribution. Pooled over 180 sequences, the median $p(\text{GT})$ at position 0 is 9.5×10^{-15} (text-bearing) and 3.7×10^{-13} (blank); the arithmetic means are 2.67×10^{-3} and 1.51×10^{-3} , close to uniform, but they are set by a handful of sequences. Only 2.8% of text-bearing sequences (5 of 180) and 3.3% of blank sequences (6 of 180) assign the correct first symbol more probability than uniform.

Table 8: Histogram of $\log p(\text{GT})$ at position 0, pooled over three seeds.

Bin ($\log p$)	Text-bear. count	%	Blank count	%
< -25	126	70.0%	115	63.9%
$[-25, -20)$	24	13.3%	25	13.9%
$[-20, -15)$	14	7.8%	13	7.2%
$[-15, -10)$	8	4.4%	13	7.2%
$[-10, -5)$	6	3.3%	10	5.6%
$[-5, 0]$	2	1.1%	4	2.2%

Seventy percent of text-bearing sequences place $\log p(\text{GT}) < -25$ at position 0. The collapse is a property of the distribution, not of its summary.

Rank of the ground truth (E1). Table 9 gives the full per-seed and pooled rank statistics referenced in Section 8. Rank is computed over the 367-symbol vocabulary at position 0 under teacher forcing, on text-bearing input, $n = 60$ per seed.

Table 9: Rank of the ground-truth symbol at position 0, text-bearing input. $\bar{H}$ is mean predictive entropy in nats.

	Median rank	IQR	$P(\text{rank} > 100)$	$\bar{H}$
Seed 0	60.5	[28.2, 104.5]	0.267	0.492
Seed 1	59.5	[28.2, 129.2]	0.333	0.285
Seed 2	95.0	[54.5, 176.5]	0.450	0.222
Pooled	70.5	[33.0, 136.2]	0.350	0.333

C Cross-Attention Contribution Norms

Full per-layer, per-seed values for $\|\text{cross-attn output}\|/\|\text{residual stream}\|$ referenced in Section 7, computed over text-bearing and blank input, $n = 60$ images per seed.

Table 10: Cross-attention contribution norm by decoder layer (L0–L4), seed, and condition.

Seed, condition	L0	L1	L2	L3	L4
0, blank	0.1131	0.1232	0.1247	0.1188	0.1068
0, text-bearing	0.1202	0.1306	0.1273	0.1191	0.1095
1, blank	0.1030	0.1172	0.1183	0.1142	0.1026
1, text-bearing	0.1269	0.1427	0.1427	0.1347	0.1249
2, blank	0.1025	0.1141	0.1217	0.1100	0.1047
2, text-bearing	0.1214	0.1316	0.1337	0.1234	0.1156

Text-bearing exceeds blank at every one of the 15 (seed, layer) cells, though the margin is modest (0.007–0.020 in absolute norm) relative to the total contribution (0.10–0.14).

D Paired Tests and Bootstrap Intervals

Because the same 60 image identifiers appear under every condition, we use paired rather than unpaired tests. Per-seed Wilcoxon signed-rank on grapheme CER (text-bearing minus blank) gives $+0.180$ ($p = 1.4 \times 10^{-6}$), $+0.088$ ($p = 0.106$) and -0.159 ($p = 0.015$) at seeds 0, 1 and 2. The sign reverses at seed 2. Pooled unclustered the difference is $+0.036$ ($p = 0.025$); clustered by seed it is not significant ($t = 0.257$, $p = 0.822$). For mean confidence the pooled difference is $+0.0016$ ($t = 0.505$, $p = 0.664$).

Seed-clustered bootstrap intervals (2,000 replicates, images resampled within seed): grapheme CER 0.9855 [0.9406, 1.0347] text-bearing and 0.9491 [0.9053, 0.9911] blank; mean confidence 0.9873 [0.9852, 0.9892] text-bearing, 0.9857 [0.9838, 0.9874] blank, 0.9857 [0.9843, 0.9871] Ol Chiki and 0.9894 [0.9871, 0.9916] Perso-Arabic; position-0 geometric mean $p(\text{GT})$ 2.20×10^{-11} [1.04×10^{-11} , 5.29×10^{-11}] text-bearing; AUROC 0.838 [0.771, 0.902] on synthetic natural and 0.570 [0.491, 0.650] on the held-out images.

E Training and Evaluation Configuration

Architecture. Vision Transformer encoder, 6 layers, $d_{\text{model}} = 320$, 5 heads, MLP ratio 4, dropout 0.1, sinusoidal position encodings, single grayscale input channel. Autoregressive decoder, 5 layers, $d_{\text{model}} = 384$, 6 heads, MLP ratio 4, dropout 0.1, learned position encodings, embedding and output head tied. Patch size 14. Total parameters 19,607,104 at $|V| = 367$. The vocabulary is built from the run’s own manifest with a minimum grapheme frequency of 5.

Training data. Rendered Hindi line crops in the natural glyph-frequency mode, 2,538 lines; the flattened and inverted modes contain 2,707 and 2,872 lines. Source text is ground-truth strings from GlotOCR Hindi [Kargaran et al., 2026], resampled into synthetic pages. Fonts are selected as the first available among Kohinor, Sangam, ITF or Noto Sans Devanagari; the specific font file used for the training run was not recorded. No augmentation is applied at training time: crops are loaded grayscale, divided by 255 and white-padded. Line crops are rendered at height 70 px (five patches) with variable width padded per batch to a multiple of the patch size.

Optimisation. AdamW, learning rate 3×10^{-4} held constant with no schedule, batch size 32, 5,000 steps, fp16 mixed precision. The checkpoint used throughout is the final step of the run; there is no validation-loss selection. Three seeds were trained under identical settings. Wall-clock was not recorded.

Evaluation data. Evaluation images are GlotOCR Bench renders, not scans, in that benchmark’s two released rendering profiles, PLAIN and OLD_DOCUMENT; the latter simulates an aged document through paper texture, ink wear, geometric distortion, downsampling and compression artifacts. They are resized to height 70 preserving aspect ratio (Lanczos). The model-side evaluation uses 60 Hindi rows in both rendering variants, 100 Ol Chiki rows and 20 Perso-Arabic rows, each under three seeds. Grapheme-cluster segmentation for every CER in this paper uses the regex library’s `\X` operator.

Scope of the engine study. The counts in Table 1 are prediction rows pooled over three engines, four languages (Hindi, Bengali, Santhali, Kashmiri) and two image variants, scored on full strings. They are not the same evaluation set as the model-side results and the overlap between the two was not recorded.

Train/evaluation overlap. `train.py` does not hold out the 60 evaluation strings. `LineDataset` appends every JSONL row of the supplied manifest, the vocabulary is built from every text field of that same file, and the documented Hindi run passes `data/manifests/hindi_natural.jsonl` (2,538 lines) with no filter. The n -gram references in Section 8 are the only fit that drops evaluation-matching lines: a manifest row is kept unless its text equals one of the 60 evaluation ground-truth strings, which leaves 2,491 lines. That is not $2,538 - 60 = 2,478$. Only 19 of the 60 evaluation strings occur as a manifest line, and those 19 occupy 47 rows because some strings repeat; $2,538 - 47 = 2,491$. The other 41 evaluation strings have no exact line match, so they are not among the dropped rows.

F Sequence Lengths

Reference transcriptions average roughly 40 graphemes (2,397 tokens over 60 strings, from Table 6), so the instrument under-generates by about 30% in both conditions. Mean and median generated grapheme lengths are 27.96 / 28.0 (text-bearing Hindi), 31.39 / 29.0 (blank), 28.28 / 29.0 (Ol Chiki) and 33.40 / 35.0 (Perso-Arabic), n as in Table 2. Blank generations are not shorter than text-bearing ones, which rules out truncation as the explanation for the CER ordering in Section 5 and leaves repetition, documented in Section 6.2, as the operative mechanism.